

The History and Future of Core Dumps in FreeBSD

Sam W. Gwydir, Texas A&M University sam@samgwydir.com

March 25, 2017

Who Am I?

My name is Sam Gwydir

- ▶ Texas A&M University - Graduating in May
- ▶ Computer Engineering/Computer Science & Mathematics
- ▶ I've used *NIX for about 12 years
- ▶ OpenBSD, and later FreeBSD for about 4


Overview

- ▶ ~~Who Am I?~~
- ▶ Background
 - ▶ What is a Coredump?
 - ▶ Procedure
 - ▶ Tutorial
- ▶ The Past
 - ▶ Timelines
- ▶ The Present
 - ▶ FreeBSD
 - ▶ Illumos
 - ▶ Mac OS X
 - ▶ Backtrace.io
- ▶ The Future
 - ▶ Core Dump Extensions
 - ▶ Modular Dump Code
 - ▶ netdump
 - ▶ minidumpsz
 - ▶ Compressed Dump

What is a Core Dump?

Core Dump A machine readable form of the state of a machine at some point in time, usually after a `panic(9)`.
Useful for debugging both userspace programs and the kernel – we will focus on the kernel today.
Written using various methods to various over the years

General Dump Procedure (4.1 BSD - FreeBSD CURRENT)



- ▶ Started by a panic(9), reboot -d
 - ▶ `sysctl debug.kdb.panic=1`
 - ▶ `dtrace -w -n 'BEGIN{ panic();}'`
- ▶ `dumpsys()` lands all/part of memory on swap in a particular format
- ▶ On reboot, `savecore(8)` writes dump to `dumpdir` for analysis

How to take a Core Dump in FreeBSD

- ▶ You are purposely panicking your machine.
- ▶ Do this in a VM!

```
root@:~ # sysrc dumpdev="AUTO" dumpdir="/var/crash"
root@:~ # mkdir /var/crash # create the dumpdir
root@:~ # chmod 700 /var/crash # fix permissions
root@:~ # # for text dump: sysctl debug.ddb.textdump.pending=1
root@:~ # sysctl debug.kdb.panic=1
```

Overview

- ▶ ~~Who Am I?~~
- ▶ ~~Background~~
- ▶ The Past
 - ▶ Timelines
- ▶ The Present
 - ▶ FreeBSD
 - ▶ Illumos
 - ▶ Mac OS X
 - ▶ Backtrace.io
- ▶ The Future
 - ▶ Core Dump Extensions
 - ▶ Modular Dump Code
 - ▶ netdump
 - ▶ minidumpsz
 - ▶ Compressed Dump

The History

- ▶ The Odyssey of `doadump()`
- ▶ Starts at 6th Edition Research UNIX's `crash(8)`
- ▶ Ends at FreeBSD 12-CURRENT's Encrypted Dump
- ▶ Turn to the Appendix for a more in depth history
 - ▶ Includes architecture support
 - ▶ Feature changes and larger bug fixes
- ▶ For even more depth, go to the org-mode file on github.
 - ▶ Includes commits, mailing list emails, copious notes.

Core Dump Output Format Time Line

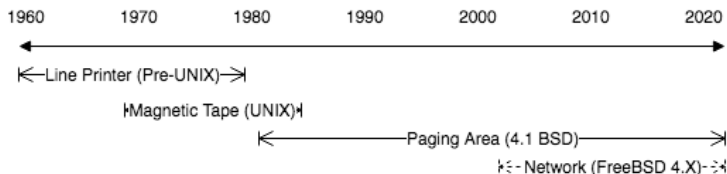


Figure 1: Core Dump Types

"Well in 1979 I can remember doing a crash dump on a Harris S/210 24 bit machine to the line printer in octal, it only took 2 hours to print. . . ." - rgrimes

FreeBSD Core Dump Extension

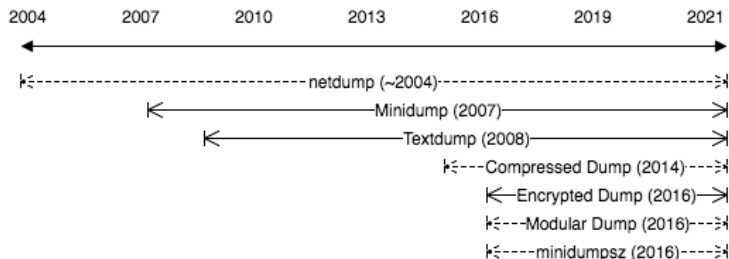


Figure 2: Core Dump Extension Timeline

Overview

- ▶ ~~Who Am I?~~
- ▶ ~~Background~~
- ▶ ~~The Past~~
- ▶ The Present
 - ▶ FreeBSD
 - ▶ Illumos
 - ▶ Mac OS X
 - ▶ Backtrace.io
- ▶ The Future
 - ▶ Core Dump Extensions
 - ▶ Modular Dump Code
 - ▶ netdump
 - ▶ minidumpsz
 - ▶ Compressed Dump

What is a Core Dump?

Two Types:

- Full Dump** Full contents of memory. Current procedure dates to 4.1 BSD. If you need help with your PDP-11 instructions are in the appendix.
- Minidump** Only active kernel pages are dumped. Added By Peter Wemm in FreeBSD 6.2. The default as of FreeBSD 7.0.

What is inside?

<code>info</code>	Metadata about dump (time, panic string, hostname)
<code>core.txt</code>	System info (backtrace, ps, vmstat, netstat, fstat)
<code>vmcore</code>	core itself

Full Dump On-Disk Format

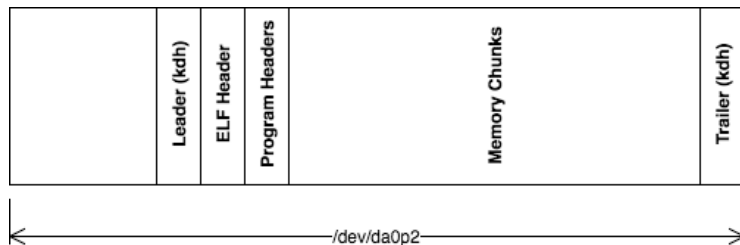


Figure 3: Full Dump Format

- ▶ Full Dump (FreeBSD 6.0)
 - ▶ A classic core dump – the full contents of memory at the time of a crash
 - ▶ ELF Format (a.out previous to FreeBSD 6.0)

Minidump On-Disk Format



Figure 4: Mini Dump Format

- ▶ Minidump (FreeBSD 6.2) - Peter Wemm
 - ▶ Contains only memory pages in use by kernel
 - ▶ Much smaller than the full contents of memory, modern dumps can still be fairly large
 - ▶ Custom “minidump” Format

What is a Text Dump?

`textdump(4)` “The textdump facility allows the capture of kernel debugging information to disk in a human-readable rather than the machine-readable form normally used with kernel memory dumps and minidumps.”

Added by Robert Watson in FreeBSD 7.1.

What is inside?

- ▶ Other files are easily added with a small patch

<code>version.txt</code>	Kernel version string
<code>panic.txt</code>	Kernel panic message
<code>msgbuf.txt</code>	Kernel message buffer
<code>config.txt</code>	Kernel configuration
<code>ddb.txt</code>	Captured DDB output

Text Dump On-Disk Format



Figure 5: Text Dump Format

- ▶ Text Dump (FreeBSD 7.1) - Robert Watson
 - ▶ Custom ddb scripting in lieu of a dump
 - ▶ Written backwards because size is unknown a priori.
 - ▶ USTAR format

Core Dumps vs Textdumps

Both

- ▶ Useful when crashes aren't predicted i.e. production
- ▶ Operators can debug crashes offline
- ▶ Allows archiving of crash data for later comparison

Core Dumps

- ▶ Do not need to know what you are looking for ahead of time
- ▶ Need source tree, debug symbols and built kernel for analysis

Text Dumps

- ▶ Less Complete but much smaller (A few MB vs Many GB)
- ▶ Sometimes easier to extract information using DDB over kgdb.

Illumos

Not a BSD but the features are alluring

- ▶ Online dump size estimation
 - ▶ Includes different calculations for settings, e.g. compression
- ▶ Compressed Dump
 - ▶ gzip compression
- ▶ Dump to Swap on zvol
 - ▶ Versatility of zvols vs partitions
- ▶ Live Dump
 - ▶ Useful for production machines where interactive debugging is not possible
 - ▶ Especially for debugging hangs

Mac OS X

- ▶ Very different from the BSD dump procedure
 - ▶ Mach-O
 - ▶ Local or remote (network or Firewire)
- ▶ netdump - kdumpd(8)
 - ▶ Using a modified tftpd(8) from FreeBSD!
- ▶ Compressed Dump
 - ▶ gzip compression
 - ▶ Both local and using kdumpd(8).
- ▶ Full Procedure in paper

Backtrace.io

- ▶ Backtrace.io curates kernel and userspace cores
- ▶ Snapshots allow for debugging on a laptop instead on crashed machine
 - ▶ Snapshots use automation to choose relevant sections of dump
- ▶ Allows for asking questions like:
 - ▶ Which panic is most common?
 - ▶ Correlated by datacenter? Storage Controller? Hard Drive Model? Timestamp?

Overview

- ▶ ~~Who Am I?~~
- ▶ ~~Background~~
- ▶ ~~The Past~~
- ▶ ~~The Present~~
- ▶ The Future
 - ▶ Core Dump Extensions
 - ▶ Modular Dump Code
 - ▶ netdump
 - ▶ minidumpsz
 - ▶ Compressed Dump

Extant Core Dump Extensions


- ▶ Modular Dump Code
 - ▶ Embedded systems
 - ▶ Mix and match features (compression but no netdump)
 - ▶ rgrimes@ for info
- ▶ netdump
 - ▶ Started at Duke by Darrell Anderson
 - ▶ Holding on since FreeBSD 4.x (~2004)
 - ▶ Continued at Sandvine and then Isilon
 - ▶ Almost part of FreeBSD 9.0
 - ▶ markj@ for info
- ▶ minidumpsz - minidump size estimation
- ▶ Compressed Dump

minidumpsz

```
% sudo sh minidumpsize_10.1.sh  
debug.mini_dump_size: 138127282176 # 138.1 GB(!!!!)
```

- ▶ Dump Size Estimation
- ▶ A “no op” version of the minidump code as a kernel module
 - ▶ minidumpsz for FreeBSD 10 and 11
 - ▶ Should be upstreamed soon
 - ▶ Email for binary (rgrimes@ or sam@samgwydir.com)

Compressed Dump



- ▶ “Save Compression”
 - ▶ gzip dump on the fly before landing in swap
- ▶ Compression Ratio 6:1 to 14:1
- ▶ A 32 GB Core becomes 5.34 GB!
- ▶ Fixing the patch so it applies to FreeBSD 12 after encrypted dump will take some work

Proposed Core Dump Extensions

- ▶ Dump to swap on zvol
 - ▶ gibbs@ offered to mentor me
- ▶ Live Dump
 - ▶ Gauge interest

How to use the appendix for Research

- ▶ Use the org-mode file
 - ▶ Includes many of the commit messages, emails, and code referenced
 - ▶ Bonus email: jkh@ calling this topic esoteric :)
 - ▶ Includes information on versions not referenced in paper
 - ▶ UNIX v5 and other incomplete sections
 - ▶ Includes notes of various levels of detail
 - ▶ Code is often included where applicable

Links

- ▶ Thanks to
 - ▶ Deb Goodkin for bringing me into the FreeBSD Community
 - ▶ Rodney Grimes for help reading PDP-11 assembly
 - ▶ Michael Dexter for coming up with this idea and for asking me to thank him.
 - ▶ You for coming!
- ▶ github.com/gwydirsam/bsd-coredump-history
- ▶ github.com/dspinellis/unix-history-repo
- ▶ people.freebsd.org/~rgrimes/index.html#kerneldump